

Workwear Site — Complete Restructure & Rebuild Plan

Status: Draft plan for review **Author:** Prepared for Damien, Phoenix Cleaning Company **Date:** 2026-07-07 **Goal:** Rebuild a large (thousands → tens of thousands of pages) non-ecommerce workwear content site that safely funnels visitors to the ecommerce store, **without** duplicate/thin/ unhelpful-content risk to either domain.

Note on scope: This repo currently contains only a README — the existing workwear build is not in it. This document is therefore a strategy + architecture plan written from your description. Where it assumes a tech stack or data source, that's flagged as an **[ASSUMPTION]** so we can correct it once the current build is available.

0. TL;DR — the one idea that fixes all three problems

The current site failed because pages were generated as **prose from a template**: same skeleton, lightly reworded body copy, keyword swapped per batch. Google's systems now explicitly target exactly this pattern ("scaled content abuse", March 2024 spam policy) regardless of whether a human or AI wrote it.

The fix is to **stop generating prose and start rendering data**. Every page must be built from a **structured dataset** (real standards, real product specs, real role/hazard requirements), so that:

- **Uniqueness is a by-product of the data**, not something we try to "spin" into copy.
- **Structure varies because the data varies** — a page about a flame-resistance standard genuinely looks different from a garment-care page, because it *contains different things*.
- **Every page earns its existence** by adding information a searcher can't easily get elsewhere ("information gain").

You *can* build tens of thousands of pages safely. It is not about the number — it's about the **value density** (average usefulness per page) and never publishing a page that has no reason to exist.

1. Diagnosis — name the actual risk, so we fix the right thing

There is no generic "duplicate content penalty." The real exposure is four specific mechanisms:

#	Mechanism	What it is	How the current build trips it
1	Scaled content abuse (spam policy, Mar 2024)	Producing many pages <i>primarily to rank</i> , with little value per page — AI or human, doesn't matter.	Batches of 300/500/1000 templated pages with reworded copy. This is the textbook definition.
2	Site-level quality classifier (helpful-content signals, now part of core ranking)	Google scores the <i>whole site's</i> average quality. A large body of thin pages drags down even the good pages.	Thousands of near-duplicate pages set a low site-wide baseline.
3	Doorway pages policy	Pages that exist mainly to funnel users onward to another destination, without standalone value.	A non-ecommerce site whose purpose is to push visitors to an ecommerce store is <i>structurally</i> at risk here — this is the single biggest cross-site danger.
4	Duplicate clustering / crawl-budget waste	Google canonicalises near-duplicates together; only one ranks, the rest waste crawl budget and dilute signals.	Batch pages that differ only by keyword collapse into one cluster.

Cross-site (the "penalise both sites" fear): the risk is *not* that linking transmits a virus. It's that (a) a site judged as a **doorway/low-value funnel** has its outbound links devalued and its referrals discounted, and (b) if the workwear site is publicly associated with the brand, its low quality is a **reputational/E-E-A-T** drag. The cleaning-services site "exposed" this because it cleared the bar the workwear site doesn't.

Design/branding is also a ranking-adjacent signal. Inconsistent, un-branded pages that "don't look like real website pages" hurt **Trust** (the T in E-E-A-T), increase pogo-sticking/bounce, and read to a quality rater as low-effort. Branding is not cosmetic here — it's part of the quality story.

2. Governing principle — the "unit of unique value" test

No page ships unless it passes at least one of these three tests. This single gate prevents the combinatorial explosion of near-duplicates.

1. **Unique information** — contains facts, data, specs, or synthesis a searcher can't trivially get from the top results already (real standard clauses, real GSM/fabric specs, real role hazards).
2. **Unique utility** — *does something*: a selector, calculator, checklist generator, comparison engine, size tool. Interactivity is inherently non-duplicable.
3. **Unique intent** — serves a genuinely distinct search need, not a synonym of an existing page. ("Class 2 vs Class 3 hi-vis" ≠ "high-visibility clothing classes" — decide which one page owns that intent and consolidate the rest.)

If a proposed page passes *none* of these, **it must not be built** — this is how we get to tens of thousands of *safe* pages instead of thousands of risky ones.

3. Content architecture — where genuine uniqueness comes from at scale

The engine of the rebuild is a **structured dataset**, not a copy template. Model workwear as entities and attributes, then *render* pages from them.

3.1 The data model (entities)

- **Garments:** hi-vis vest, coverall, work trouser, FR jacket, safety boot, chef whites, scrub, lab coat, tabard, softshell, waterproof, etc.
- **Sectors:** construction, healthcare/NHS, catering/hospitality, warehousing/logistics, manufacturing, automotive, electrical, rail, highways, agriculture, cleaning, security.
- **Roles:** electrician, welder, chef, warehouse operative, mechanic, scaffolder, groundworker, nurse...
- **Standards/certifications (the richest unique-data seam):** EN ISO 20471 (hi-vis), EN ISO 11612 (heat/flame), EN 343 (rain), EN ISO 20345 (safety footwear), IEC 61482 (arc flash), EN 1149 (antistatic), EN 13034 (chemical), rail GO/RT, etc.
- **Attributes:** fabric, GSM/weight, colour, reflective-tape class, features (knee-pad pockets), care/wash temperature, sizing.
- **Intents:** buying guide, "what to wear for X", compliance explainer, care/maintenance, sizing, comparison, cost, branding/embroidery.

3.2 The unique-data backbone: your ecommerce catalogue [ASSUMPTION: catalogue data is accessible]

Your ecommerce product catalogue is the asset the competition can't copy: **real** fabric, GSM, certifications, colours, price bands, sizing. Spec-driven pages built from it are:

- Genuinely unique (real numbers in real tables, not prose).
- Naturally, *relevantly* linked to the store — which is the **antidote to the doorway problem** (the link is a citation of real data, not a funnel CTA).
- Cheap to keep fresh (regenerate when the catalogue changes).

3.3 Discipline against combinatorial near-duplicates

`sector × garment × role × region` is millions of combinations — and 95% of them would be near-duplicates. Rules:

- **Only materialise a combination when the data differs meaningfully.** "FR workwear for welders" is real (arc flash + heat + spatter). "Hi-vis vests in Swindon" is a doorway — **do not build pure geo-programmatic pages** unless backed by real local data (stockists, local regs, delivery).
 - **Variable depth:** high-value nodes become full pages; thin nodes become a *section within* a parent page or a filtered view, never a standalone URL.
 - **Cap each series by value, not by a target count.** Kill the "build 1000" mindset — build "build every node that passes §2, and no more."
-

4. Page taxonomy & template families (fixing the "not enough structural variation" problem)

The old build used ~1 template per series. The rebuild uses **many distinct template families**, each with genuinely different DOM/structure because each renders different data:

Page type	Purpose	Why it's structurally distinct	Rough scale
Standard/regulation explainer	Explain EN ISO 20471 etc.	Clause tables, class diagrams, "does this apply to me" logic	1 per standard (dozens)
Sector guide (hub)	Workwear for construction, NHS...	Hazard matrix, role list, required standards, seasonal	1 per sector (~15–25)
Role guide	What a welder/electrician wears	Hazard → garment → standard mapping, layering	1 per meaningful role (100s)
Garment spec/comparison	Polycotton vs cotton; Class 2 vs 3	Spec tables, decision tables, pros/cons from real data	100s
Care & maintenance	Wash FR coveralls without losing FR	Step lists, temp/care tables, do/don't — high utility	100s
Sizing & fit	Size guides per garment/brand	Measurement tables, fit diagrams	100s
Interactive tools	Hi-vis class selector, "which coverall", PPE checklist, embroidery cost, size calculator	Pure JS utility — inherently unique	5–15 flagship tools
Glossary/definitions	Define terms/standards/fabrics	Short, interlinked, schema-marked	100s
Editorial/blog	News, seasonal, deep how-tos, case studies	Human-led, opinion/experience (E-E-A-T)	ongoing

Structural-variation rules baked into the generator:

- Each family has **3–5 layout variants**, and block order/inclusion is **driven by which data fields are present**, so no two pages in a series share an identical skeleton.
- A **component library** (spec table, hazard matrix, comparison grid, callout, FAQ, tool embed, related-links) assembled by rules — not a fixed sequence.
- **Machine-similarity ceiling**: the QA gate (§6) rejects any page whose structure/copy exceeds a similarity threshold vs its siblings.

5. Branding & design system (fixing "doesn't look like a real website")

Build **one** design system before any page is generated. Every page renders through it, so consistency is automatic, not manual.

- **Design tokens**: a defined colour palette (primary/secondary/neutral/semantic), type scale, spacing scale, radii, shadows — in one tokens file.
- **Global chrome**: one header, nav, footer, breadcrumb pattern on every page.

- **Component library:** buttons, cards, tables, callouts, tool shells — all themed from tokens.
- **Imagery/iconography system:** consistent treatment; real photography where possible over generic stock (helps E-E-A-T and originality).
- **Accessibility baked in** (contrast, focus states, semantic HTML) — also a quality signal.
- **Deliverable:** a short brand/design spec + a live component/style guide page.

If you have brand assets from the ecommerce store, we mirror its palette/type so the two sites feel like one family — reinforcing trust across the funnel.

6. Build system & quality gates (the part that makes scale safe)

6.1 Recommended stack [ASSUMPTION — confirm/replace]

- **Static-site generator** (Astro or Next.js) — fast, indexable, component-driven.
- **Content as data:** the entity dataset in structured files/DB; pages rendered from it via the component library and template families. Editorial in MDX/CMS.
- **Clean separation:** data/ (facts) — templates/ (families) — components/ (blocks) — content/ (editorial). Copy is never hand-duplicated across pages.

6.2 Pre-publish QA gates (automated — nothing ships that fails)

1. **Uniqueness / similarity check** — embeddings or shingling (e.g. SimHash/MinHash) to compare each page against its siblings; **reject** above a similarity threshold.
2. **Thin-content check** — minimum genuine information density (not just word count): must contain real data blocks (spec table, hazard matrix, tool, etc.), not just prose.
3. **\$2 value-test gate** — each page tagged with which of the three tests it passes; untagged = blocked.
4. **Doorway check** — CTA-to-content ratio ceiling; a page can't be mostly "buy now" links.
5. **Technical SEO** — unique title/meta/H1, correct canonical, schema.org markup (Product, FAQ, HowTo, Definition as appropriate), internal links present, no orphan pages.
6. **Design/lint** — passes design-system + accessibility + Core Web Vitals budget.

6.3 Index management

- **Phased publishing** — release in tranches; watch indexation and rankings before the next tranche.
- **noindex until proven** for experimental series; promote to index only once quality is confirmed.
- **XML sitemaps segmented by page type** so we can monitor how Google treats each family.

7. Dealing with the existing pages (don't just leave them)

The old pages are actively dragging the site-wide score. Audit and triage **every** URL:

1. **Crawl + score** existing pages: traffic, impressions, similarity cluster, thinness.
2. **Triage into four buckets:**

- **Keep & upgrade** — has intent value; rebuild on the new data model.
- **Consolidate** — merge near-duplicate clusters into one strong canonical page; 301 the rest.
- **Prune** — no value, no traffic, no unique intent → **410/redirect**. Removing thin pages *raises* the site-wide average; this is a known recovery lever.
- **Rewrite as editorial** — salvage genuinely useful topics as human-led content.

3. **Migration map** with 301s so equity is preserved and no soft-404s remain.

Expect to *delete a large share* of the current pages. That is a feature, not a loss.

8. Cross-site linking strategy (protect both domains)

- **Every outbound link to the store is a relevant citation** ("this coverall meets EN ISO 11612 — view specs/buy"), not a blanket funnel CTA. Relevance is what separates a resource from a doorway.
- **The workwear site must stand alone** — genuinely useful with zero purchases made. If it passes §2 on its own merits, it is not a doorway by definition.
- **Reasonable link density** — content first, commercial links secondary and contextual.
- **Consistent brand association** so trust flows both ways once quality is established.
- Follow links to your own store are fine; the protection is the *quality of the linking page*, not link attributes.

9. Metrics & governance

- **Value density** — % of indexed pages passing §2 (target: 100%; this is the north star).
- **Indexation ratio** — indexed / submitted per page-type sitemap (near-duplicate families show low ratios).
- **Similarity distribution** — internal duplicate-cluster report, tracked over time.
- **Organic performance** — impressions/clicks per page family; funnel referrals to the store.
- **Core Web Vitals + design conformance** — automated budgets.
- **Editorial calendar** for the human-led layer (ongoing freshness + E-E-A-T).

10. Phased roadmap

Phase	Outcome	Key deliverables
0. Audit & decisions	Know the current state; lock the stack & data sources	URL inventory + triage buckets; confirm catalogue data access; confirm SSG stack
1. Foundations	The "one site" exists	Design system + tokens + component library; brand spec; base SSG scaffold
2. Data model	Facts, not prose	Entity dataset; standards/spec data ingested from catalogue; §2 tagging schema
3. Template families	Structural variety	3–5 layout variants per page type; component-assembly rules
4. QA harness	Scale becomes safe	Similarity, thin-content, doorway, technical-SEO, design gates in CI
5. Pilot tranche	Prove the model	Build 1–2 high-value families (e.g. standards explainers + sector guides), index, monitor
6. Old-page cleanup	Raise site-wide average	Execute consolidate/prune/redirect migration map
7. Scale out	Thousands → tens of thousands, safely	Roll out remaining families in monitored tranches
8. Editorial + tools	E-E-A-T + utility moat	Flagship interactive tools; ongoing editorial calendar

11. Open questions to confirm before Phase 1

1. **Where is the current build**, and what stack was it generated in? (Determines migration effort.)
2. **Can we access the ecommerce product catalogue as structured data?** (This is the uniqueness engine.)
3. **Preferred stack** — is Astro/Next acceptable, or is there an existing platform to stay on?
4. **Brand assets** — do we have the ecommerce store's palette/type/logo to mirror?
5. **URL/domain** — same domain rebuild, or new structure with redirects?
6. **Scale ambition & timeline** — realistic first-tranche size and go-live target?

Appendix A — Worked example: one node, done right

Topic: "Flame-resistant coveralls for welders"

- **Passes §2:** unique information (welding hazards: arc, spatter, radiant heat → EN ISO 11612 + IEC 61482 requirements) **and** unique utility (embedded "FR standard finder" tool).
- **Data rendered:** required standards table, fabric/GSM options from catalogue, care rules that preserve FR properties, layering guidance, sizing.

- **Structure:** hazard matrix → standards table → real product specs → care do/don't → tool → FAQ.
Different DOM from a sizing page because it contains different data.
- **Outbound link:** contextual — "coveralls certified to EN ISO 11612 in the catalogue" (citation, not funnel).
- **Result:** genuinely useful with zero purchase; not thin; not duplicate; not a doorway.

Contrast the old build: a template that said "Looking for FR coveralls for [ROLE]? We have a great range..." with the role swapped — thin, duplicative, doorway. **That page type is retired.**