

iNeedWorkwear — Complete Restructure & Rebuild Plan (v2)

Status: Plan for review — now grounded in the actual build **Prepared for:** Damien, Phoenix Cleaning Company **Date:** 2026-07-07 **Goal:** Rebuild the iNeedWorkwear content site so it can safely reach thousands → tens of thousands of genuinely useful pages that funnel to the (not-yet-built) ecommerce store, without duplicate/thin/doorway-content risk to *either* domain.

v2 note: This version replaces the earlier assumptions-based draft. I've now reviewed a 400-file sample of the real workwear build, the [CW Hybrid Series Planning Matrix](#), and four pages from the cleaning site (Phoenix Duct Clean) that you hold up as the better build. The findings below are measured from those files, not assumed.

0. Headline recommendation: start fresh — but harvest three assets

You asked whether to rebuild, edit, or scrap. **Scrap the current workwear pages as a live site.** They are a net liability, not a foundation — and editing 10,000 pages to fix problems baked into the template and the topic list costs more than rebuilding and leaves the risk signals in place.

But "start fresh" does **not** mean start from nothing. Harvest three things from what exists:

1. **The design-system discipline of the cleaning site.** Phoenix Duct Clean already proves your team can build one coherent, well-branded system (one palette, one type system, schema markup, consistent chrome). That *is* the model. The workwear site's core failure is that it never had one.
2. **The industry taxonomy** in the planning matrix (construction, healthcare, automotive, plumbing... with product mappings). The *entity model* is sound; the *town-multiplication* on top of it is not.
3. **A handful of genuine editorial guides** (e.g. "best safety boots for concrete floors"). These can be rebuilt as real, deep buying guides — the seed of the useful-content layer.

Everything else — the town-swapped industry pages and, urgently, the off-topic traffic-net pages — should not be migrated. Redirect or [410](#) them (see §7).

1. What the actual build shows (evidence)

1.1 Branding — no system at all (your problem #1, confirmed)

The cleaning site uses **one** design-token block on every page: `--steel:#0E1620; --amber:#D4762A; --teal:#3FA89B`, fonts `Saira Condensed / Inter / Space Mono`. Every page is unmistakably the same site.

The workwear site has **a different palette and often different fonts per series**:

Series	Sample colour	Fonts
am- (automotive)	#1c2024 dark + amber	Lora / Source Sans
bh- (hair & beauty)	#271f33 purple	Lora / Source Sans
festival-	#7c3aed violet	Lora / Source Sans
restaurant-	#2c2a27	Playfair Display / DM Sans
blog-	#063a09 green	(varies)
eat-	#13251a green	Lora / Source Sans

There is no shared header/nav, no shared footer identity, no tokens. Each batch looks like a different website. This is exactly the "doesn't look like real website pages" problem — and it's the single clearest difference from the cleaning site.

1.2 Uniqueness — spun content around a fixed skeleton (your problem #2, confirmed)

Measured on two pages from the same series (am-aberdeen vs am-abingdon , both "Garage and Mechanic Workwear Supplier"):

- **77% shared vocabulary** (2,216 of ~2,880 word tokens common to both).
- **7 of 8 <h2> headings identical, word-for-word.** Only the opening "local flavour" heading differs ("The energy city's fleet and 4x4 trade" vs "The MG factory and the Radley Road trade").
- Body sentences reworded around the same fixed structure.

That is the classic **spun-content signature**: same skeleton, same topic, lightly reworded prose, one town swapped in. Google's systems detect this as near-duplication even when exact phrase overlap is low, and cluster the pages together.

1.3 Templating — duplication by design (your problem #3, confirmed)

The CW Hybrid Series Planning Matrix sets out the intended system:

- **~37 industry series**, most at **500 towns** each (some 200), tiered for build order.
- Rendered through **two templates** — "A (Concierge)" and "B (Self-checkout)" — plus a hybrid.
- The declared uniqueness lever is a "Local Variation: High/Medium/Low" column — i.e. **town-name substitution**. "Towns: 500" against each industry.

That's **~10,000+ near-identical pages across 2–3 templates**. The matrix is, in effect, a blueprint for the exact pattern ("scaled content abuse") that Google's March 2024 spam policy targets.

1.4 The fourth problem — topical sprawl / doorway pages (not on your list; the biggest risk)

The sample contains large programmatic series with **nothing to do with workwear**, each keyword-stuffed and linking to the shop:

Page	Words	"workwear" mentions	Shop links
world-cup-bulgaria (Bulgaria's WC history)	2,882	31	11
festival-2000-trees (music festival guide)	2,794	43	17
eat-barnet ("best places to eat")	2,516	45	12
best-youth-clubs-harrow	2,625	12	5
restaurant-carlton ("launch a restaurant")	3,058	21	6

Others in the set: town criers, nurseries, festivals A–Z, "eat" guides by town, world-cup by nation. A World Cup history page on a workwear domain that mentions "workwear" 31 times and pushes 11 shop links is a textbook **doorway page**. This is the content most likely to get the domain classified as a low-value funnel — and because these pages link to the ecommerce store, that judgement is what puts **both** sites at risk. **This series must be deleted, not fixed.**

1.5 A caution about the cleaning site (the "good" build)

The cleaning site is genuinely better — one design system, real local specifics (street names, councils, local history), useful editorial. **But do not simply clone its location-page pattern at 10,000× scale**, because it carries two latent risks that only bite at scale:

- **Repeated fixed "stats"** ("4,287 canopies degreased", "54,754 hours on site") appear identical on every page. If those numbers aren't literally true, that's an E-E-A-T/trust problem; at scale it's a duplication signal too.
- **"On the ground" job anecdotes and testimonials** (e.g. the "bacon rolls" story) read as per-town template slots. If they're not real, that's a trust risk on a page type Google scrutinises.

The cleaning site works because it is **small and tightly on-topic** (a few hundred pages in one niche). The workwear ambition is 20–50× larger and was topically sprawling — the same location-page recipe that passes at cleaning-site scale would trip the **site-level quality classifier** at workwear scale. That gap is the heart of your question, and §2 is the answer to it.

2. Can you build tens of thousands of pages safely? Yes — but not this way

The honest answer: **there is no safe page count — there is only safe value density**. 50,000 genuinely useful pages is fine; 5,000 town-swapped ones is not. Google scores the *average* quality of the site, so the constraint is: *every page must add something a searcher can't easily get elsewhere*.

The current build fails because its only uniqueness lever is the town name. To scale safely, uniqueness has to come from **real, structured data and real utility** — things that genuinely differ page to page. The rest of this plan is how to get there.

The governing rule — the "unit of unique value" test

No page ships unless it passes at least one of these three tests:

1. **Unique information** — real facts a searcher can't trivially get from the current top results (a specific safety standard's clauses, real garment specs, a real role's hazard profile).
2. **Unique utility** — it *does something*: a selector, calculator, checklist generator, comparison engine. Interactivity is inherently non-duplicable.
3. **Unique intent** — it serves a genuinely distinct search need, not a synonym of another page.

A town-swapped "Garage Workwear in [Town]" page passes none of these. A "Bulgaria World Cup" page passes none of these *for a workwear site*. Both are retired. This one gate is what converts a 10,000-page liability into a 10,000-page asset.

3. Where genuine uniqueness comes from (the engine)

Stop generating **prose from a template**. Start rendering **data through a component library**. Uniqueness then becomes a by-product of the data, and structure varies because the data varies.

3.1 The data model (entities), not the town list

- **Garments:** hi-vis, coverall, work trouser, FR jacket, safety boot, chef whites, scrub, tabard, softshell, waterproof...
- **Sectors:** construction, healthcare, hospitality, warehousing, automotive, plumbing/electrical, landscaping, FM/cleaning, manufacturing... (*reuse the matrix's taxonomy*)
- **Roles:** welder, electrician, chef, warehouse operative, groundworker, mechanic, nurse...
- **Standards (the richest unique-data seam):** EN ISO 20471 (hi-vis), EN ISO 11612 (heat/flame), EN 343 (rain), EN ISO 20345 (safety footwear), IEC 61482 (arc flash), EN 1149 (antistatic), EN 13034 (chemical), rail GO/RT...
- **Attributes:** fabric, GSM, colour, tape class, features, wash/care temperature, sizing.
- **Intents:** buying guide, "what to wear for X", compliance explainer, care, sizing, comparison, cost, embroidery/branding.

3.2 The supplier-DAM trap (critical — this changes the strategy)

You told me catalogue data and imagery will come from **workwear suppliers' digital asset libraries**. That means the raw specs and photos are **identical to every other reseller** pulling the same DAM — the *opposite* of a unique asset. Rules that follow directly:

- **Never publish supplier copy or specs as page text.** Treat the DAM as *raw input to transform* — synthesise your own spec tables, comparisons, and standard/hazard mappings on top of it.
- **Differentiate imagery.** Raw DAM photos will be image-deduped by Google against dozens of rivals. Plan for own photography where feasible, or at minimum processed/annotated/composited images.
- **The unique value lives in the layer you add** — standards analysis, role/hazard mapping, care guidance, tools, comparisons. No supplier feed contains those. That layer is your moat.

3.3 Discipline against combinatorial near-duplicates

sector × garment × role × town is millions of combinations, ~95% near-duplicate. So:

- **Only build a combination when the data genuinely differs.** "FR coveralls for welders" is real (arc + spatter + radiant heat). "Garage workwear in Abingdon" is not — it's a town swap.
- **Kill the town-multiplication default.** Do **not** build 500 town variants per industry. Build *one strong industry/role guide*, and create a local page **only** where real local data exists (a physical stockist, local regulation, real delivery/lead-time info) — otherwise it's a doorway.
- **Variable depth:** thin nodes become a *section within* a parent page or a filtered view, never a standalone URL.
- **Cap by value, not by a target count.** Retire "build 500" thinking. Build every node that passes §2, and no more.

4. Page taxonomy & template families (fixes the over-templating)

The old build: 2–3 templates for everything. The rebuild: **many families, each structurally distinct because each renders different data**, and each with 3–5 layout variants.

Page type	Example	Why structurally distinct
Standard/regulation explainer	"EN ISO 20471 hi-vis classes explained"	Clause tables, class diagrams, applicability logic
Sector guide (hub)	"Workwear for construction"	Hazard matrix, role list, required standards
Role guide	"What a welder should wear"	Hazard → garment → standard mapping
Garment comparison	"Class 2 vs Class 3 hi-vis"	Decision tables, spec grids from real data
Care & maintenance	"Washing FR coveralls without losing FR"	Step lists, temperature/care tables — high utility
Sizing & fit	Per-garment size guides	Measurement tables, fit diagrams
Interactive tools	Hi-vis class selector; "which coverall?"; PPE checklist; embroidery cost; size calculator	Pure utility — inherently unique
Glossary	Standards/fabric definitions	Short, interlinked, schema-marked
Editorial/blog	Buying guides, deep how-tos	Human-led, real experience (E-E-A-T)

Structural-variation rules in the generator: block order and inclusion driven by *which data fields are present*, so no two pages in a series share an identical skeleton; a component library (spec table, hazard matrix, comparison grid, callout, FAQ, tool embed) assembled by rules; and a QA gate (§6) that rejects any page too similar to its siblings.

5. Branding & design system (fixes problem #1) — systematise the cleaning site

You already have the exemplar. Do for iNeedWorkwear what Phoenix Duct Clean did — **once, centrally** — and render every page through it:

- **One token file:** palette (primary/secondary/neutral/semantic), type scale, spacing, radii, shadows. One palette for the whole site, chosen to sit alongside the future store.
 - **Global chrome:** one header, nav, footer, breadcrumb on every page.
 - **Component library** themed from tokens (buttons, cards, tables, callouts, tool shells).
 - **Imagery system** (see §3.2 — differentiate from DAM stock).
 - **Accessibility + Core Web Vitals** budgets baked in (both are quality signals).
 - **Deliverable:** a brand spec + a live style-guide page, before any content page is generated.
-

6. Build system & QA gates (what makes scale safe)

6.1 Stack [to confirm]

- **Static-site generator** (Astro or Next.js) — component-driven, fast, indexable.
- **Content as data:** the entity dataset in structured files/DB, rendered via template families and the component library; editorial in MDX/CMS. Copy is never hand-duplicated across pages.
- **Clean separation:** `data/` (facts) · `templates/` (families) · `components/` (blocks) · `content/` (editorial).

6.2 Pre-publish gates (automated in CI — nothing ships that fails)

1. **Similarity gate** — embeddings or SimHash/MinHash vs siblings; reject above a threshold. (*This alone would have blocked the `am-aberdeen` / `am-abingdon` pattern.*)
2. **Thin-content gate** — must contain real data blocks (spec table, hazard matrix, tool), not just prose.
3. **§2 value-test gate** — each page tagged with which test it passes; untagged = blocked.
4. **Doorway gate** — CTA/shop-link-to-content ratio ceiling. (*This would have blocked the World Cup pages.*)
5. **Topical-relevance gate** — page must map to a workwear entity; off-topic = blocked.
6. **Technical SEO** — unique title/meta/H1, correct canonical, schema.org, internal links, no orphans.
7. **Design/lint** — design-system + accessibility + CWV budgets.

6.3 Index management

- **Phased publishing** in tranches; watch indexation and rankings before the next tranche.
 - `noindex` until proven for experimental families.
 - **Sitemaps segmented by page type** to monitor how Google treats each family.
-

7. Dealing with the existing pages

They are actively dragging any shared quality signal. Triage every URL:

1. **Off-topic traffic-net series** (world cup, festivals, eat, youth clubs, restaurants, nurseries, town criers): **delete** → **410** (or redirect to the closest relevant hub only where one exists). Do not migrate. These are pure liability.
2. **Town-swapped industry pages: retire the town variants.** Keep the *industry topic*; rebuild as a smaller set of deep sector/role guides on the new system. 301 each dead town URL to its sector hub.
3. **Genuine editorial** (real buying guides): **rebuild** as deep, first-hand editorial. Keep the URL where it has any equity.
4. **Migration map** with 301s; no soft-404s left behind.

Expect to delete the large majority of current URLs. Removing thin/off-topic pages *raises* the site's average quality — it's a recovery lever, not a loss.

8. Cross-site linking (protect both domains)

- **Every link to the store is a relevant citation** ("coveralls certified to EN ISO 11612 → view specs"), never a blanket funnel CTA. Relevance is what separates a resource from a doorway.
- **The content site must stand alone** — useful with zero purchases. If a page passes §2 on its own merits, it is not a doorway by definition.
- **Content first, commercial links secondary** (the opposite of the current World Cup pages).
- Because the store isn't built yet, **design both around one brand system now** so they read as one family when the store launches.

9. Metrics & governance

- **Value density** — % of indexed pages passing §2 (north star; target 100%).
- **Indexation ratio** per page-type sitemap (near-duplicate families show low ratios).
- **Internal similarity distribution** — duplicate-cluster report over time.
- **Organic performance** per family + funnel referrals to the store.
- **CWV + design conformance** — automated budgets.
- **Editorial calendar** for the human-led layer (freshness + E-E-A-T).

10. Phased roadmap

Phase	Outcome	Key deliverables
0. Decommission risk	Stop the bleeding	Delete/ 410 off-topic series; 301 town pages to hubs; migration map
1. Foundations	"One site" exists	Design system + tokens + component library (systematise the cleaning site); brand spec
2. Data model	Facts, not prose	Entity dataset (reuse matrix taxonomy); standards data; DAM-transform pipeline; §2 tagging
3. Template families	Structural variety	3–5 layout variants per page type; component-assembly rules
4. QA harness	Scale becomes safe	Similarity, thin, doorway, relevance, technical-SEO, design gates in CI
5. Pilot tranche	Prove the model	Build 1–2 high-value families (standards explainers + sector guides), index, monitor
6. Scale out	Thousands, safely	Roll out remaining families in monitored tranches
7. Tools + editorial	Utility + E-E-A-T moat	Flagship interactive tools; ongoing editorial

11. Open questions before Phase 1

1. **Scope of the cull** — are you comfortable deleting the off-topic series outright (my strong recommendation), or do you want a page-by-page review first?
2. **Stack** — Astro/Next acceptable, or an existing platform (the current pages look hand-built/cPanel-hosted) to stay on?
3. **Imagery** — can we commission any own photography, or is it DAM-only at launch? (Sets how hard we must work on image differentiation.)
4. **Are the cleaning site's stats/testimonials real?** (Determines whether that pattern is safe to carry into the workwear system.)
5. **URL/domain** — rebuild on the same `iNeedWorkwear` domain with redirects, or fresh structure?
6. **First-tranche scope & timeline.**

Appendix A — Worked example: one node, done right

Topic: "Flame-resistant coveralls for welders"

- **Passes §2:** unique information (welding hazards → EN ISO 11612 + IEC 61482) **and** unique utility (embedded FR-standard finder).
- **Data rendered:** required-standards table, fabric/GSM options (transformed from DAM, not pasted), care rules that preserve FR properties, layering, sizing.

- **Structure:** hazard matrix → standards table → spec comparison → care do/don't → tool → FAQ. A different DOM from a sizing page, because it contains different data.
- **Outbound link:** contextual — "coveralls certified to EN ISO 11612 →" (citation, not funnel).
- **Result:** useful with zero purchase; not thin, not duplicate, not a doorway.

Contrast with the retired pattern: am-abingdon — "Garage and Mechanic Workwear Supplier" with the town swapped and 7/8 headings identical to am-aberdeen ; or world-cup-bulgaria — a football-history essay stuffed with 31 "workwear" mentions and 11 shop links. Both fail all three tests. **That is the build we are replacing.**