

iNeedWorkwear — Complete Restructure & Rebuild Plan (v3)

Status: Strategy + prototyping phase complete — foundations proven, awaiting the canary gate **Prepared for:** Damien, Phoenix Cleaning Company **Date:** 2026-07-07 **Goal:** Rebuild the iNeedWorkwear content site so it can safely reach thousands → tens of thousands of genuinely useful pages that funnel to the (not-yet-built) ecommerce store, without duplicate/thin/doorway-content risk to *either* domain.

v3 note: v2 was grounded in the real build (a 400-file sample, the CW Hybrid Series Planning Matrix, and cleaning-site pages). v3 folds in the **prototyping phase**: a working prototype now exists for every distinct page type, the design system is established, the industry-driven local model is validated across three towns, the hub-and-spoke content method is proven, and the pre-scale **canary gate** is defined. See §0.5 for what's built, and Appendix B for the asset index.

0. Headline recommendation: start fresh — but harvest three assets

You asked whether to rebuild, edit, or scrap. **Scrap the current workwear pages as a live site.** They are a net liability, not a foundation — and editing 10,000 pages to fix problems baked into the template and the topic list costs more than rebuilding and leaves the risk signals in place.

But "start fresh" does **not** mean start from nothing. Harvest three things from what exists:

1. **The design-system discipline of the cleaning site.** Phoenix Duct Clean already proves your team can build one coherent, well-branded system (one palette, one type system, schema markup, consistent chrome). That *is* the model. The workwear site's core failure is that it never had one.
2. **The industry taxonomy** in the planning matrix (construction, healthcare, automotive, plumbing... with product mappings). The *entity model* is sound; the *town-multiplication* on top of it is not.
3. **A handful of genuine editorial guides** (e.g. "best safety boots for concrete floors"). These can be rebuilt as real, deep buying guides — the seed of the useful-content layer.

Everything else — the town-swapped industry pages and, urgently, the off-topic traffic-net pages — should not be migrated. Redirect or 410 them (see §7).

0.5 Where we are now (v3) — foundations proven

Since v2 we've moved from strategy to **working prototypes**. The design system and the full page-type library now exist as real, rendered pages — so the remaining work is *content volume on proven templates*, not new design. What's been built and validated:

Layer	Status	Proof
Design system (iNeedWorkwear brand)	✓ Established	Navy/orange/hi-vis-lime palette, Barlow Condensed/Inter/IBM Plex Mono, one chrome — used on every prototype
Role guide + interactive tool	✓ Prototyped	fr-coveralls-for-welders (standard-finder tool, real EN ISO 11611/11612/IEC 61482 data)
Local / area (industry-driven)	✓ Prototyped ×3	Aberdeen (offshore) · Corby (logistics) · Harrogate (hospitality) — 35% shared vocab vs the old site's 77%
Content hub (pillar)	✓ Prototyped	hi-vis-workwear-guide
Standard explainer + comparison (spokes)	✓ Prototyped	en-iso-20471... + class-2-vs-class-3... , interlinked hub-and-spoke
Care / maintenance (HowTo)	✓ Prototyped	how-to-wash-fr-coveralls (numbered steps, HowTo schema)
Page architecture & counts	✓ Mapped	~1,350 safe core pages across 14 families (page-architecture.csv)
Local data engine	✓ Built	15 economic archetypes + 495-town list + local-page data template
Staging & safe launch	✓ Defined	§6.4
Pre-scale decision gate	✓ Defined	§6.5 + canary-watch-framework.md

The one thing still gating a live build is the canary (§6.5) — reading how Google treats the freshly-launched cleaning 3,000 before we replicate the model on workwear. Everything above is ready and unpublished until that reads green.

1. What the actual build shows (evidence)

1.1 Branding — no system at all (your problem #1, confirmed)

The cleaning site uses **one** design-token block on every page: `--steel:#0E1620; --amber:#D4762A; --teal:#3FA89B`, fonts `Saira Condensed / Inter / Space Mono`. Every page is unmistakably the same site.

The workwear site has **a different palette and often different fonts per series**:

Series	Sample colour	Fonts
am- (automotive)	#1c2024 dark + amber	Lora / Source Sans
bh- (hair & beauty)	#271f33 purple	Lora / Source Sans
festival-	#7c3aed violet	Lora / Source Sans
restaurant-	#2c2a27	Playfair Display / DM Sans
blog-	#063a09 green	(varies)
eat-	#13251a green	Lora / Source Sans

There is no shared header/nav, no shared footer identity, no tokens. Each batch looks like a different website. This is exactly the "doesn't look like real website pages" problem — and it's the single clearest difference from the cleaning site.

1.2 Uniqueness — spun content around a fixed skeleton (your problem #2, confirmed)

Measured on two pages from the same series (am-aberdeen vs am-abingdon , both "Garage and Mechanic Workwear Supplier"):

- **77% shared vocabulary** (2,216 of ~2,880 word tokens common to both).
- **7 of 8 <h2> headings identical, word-for-word.** Only the opening "local flavour" heading differs ("The energy city's fleet and 4x4 trade" vs "The MG factory and the Radley Road trade").
- Body sentences reworded around the same fixed structure.

That is the classic **spun-content signature**: same skeleton, same topic, lightly reworded prose, one town swapped in. Google's systems detect this as near-duplication even when exact phrase overlap is low, and cluster the pages together.

1.3 Templating — duplication by design (your problem #3, confirmed)

The CW Hybrid Series Planning Matrix sets out the intended system:

- **~37 industry series**, most at **500 towns** each (some 200), tiered for build order.
- Rendered through **two templates** — "A (Concierge)" and "B (Self-checkout)" — plus a hybrid.
- The declared uniqueness lever is a "Local Variation: High/Medium/Low" column — i.e. **town-name substitution**. "Towns: 500" against each industry.

That's **~10,000+ near-identical pages across 2–3 templates**. The matrix is, in effect, a blueprint for the exact pattern ("scaled content abuse") that Google's March 2024 spam policy targets.

1.4 The fourth problem — topical sprawl / doorway pages (not on your list; the biggest risk)

The sample contains large programmatic series with **nothing to do with workwear**, each keyword-stuffed and linking to the shop:

Page	Words	"workwear" mentions	Shop links
world-cup-bulgaria (Bulgaria's WC history)	2,882	31	11
festival-2000-trees (music festival guide)	2,794	43	17
eat-barnet ("best places to eat")	2,516	45	12
best-youth-clubs-harrow	2,625	12	5
restaurant-carlton ("launch a restaurant")	3,058	21	6

Others in the set: town criers, nurseries, festivals A–Z, "eat" guides by town, world-cup by nation. A World Cup history page on a workwear domain that mentions "workwear" 31 times and pushes 11 shop links is a textbook **doorway page**. This is the content most likely to get the domain classified as a low-value funnel — and because these pages link to the ecommerce store, that judgement is what puts **both** sites at risk. **This series must be deleted, not fixed.**

1.5 A caution about the cleaning site (the "good" build)

The cleaning site is genuinely better — one design system, real local specifics (street names, councils, local history), useful editorial. **But do not simply clone its location-page pattern at 10,000× scale**, because it carries two latent risks that only bite at scale:

- **Repeated fixed "stats"** ("4,287 canopies degreased", "54,754 hours on site") appear identical on every page. If those numbers aren't literally true, that's an E-E-A-T/trust problem; at scale it's a duplication signal too.
- **"On the ground" job anecdotes and testimonials** (e.g. the "bacon rolls" story) read as per-town template slots. If they're not real, that's a trust risk on a page type Google scrutinises.

The cleaning site works because it is **small and tightly on-topic** (a few hundred pages in one niche). The workwear ambition is 20–50× larger and was topically sprawling — the same location-page recipe that passes at cleaning-site scale would trip the **site-level quality classifier** at workwear scale. That gap is the heart of your question, and §2 is the answer to it.

2. Can you build tens of thousands of pages safely? Yes — but not this way

The honest answer: **there is no safe page count — there is only safe value density**. 50,000 genuinely useful pages is fine; 5,000 town-swapped ones is not. Google scores the *average* quality of the site, so the constraint is: *every page must add something a searcher can't easily get elsewhere*.

The current build fails because its only uniqueness lever is the town name. To scale safely, uniqueness has to come from **real, structured data and real utility** — things that genuinely differ page to page. The rest of this plan is how to get there.

The governing rule — the "unit of unique value" test

No page ships unless it passes at least one of these three tests:

1. **Unique information** — real facts a searcher can't trivially get from the current top results (a specific safety standard's clauses, real garment specs, a real role's hazard profile).
2. **Unique utility** — it *does something*: a selector, calculator, checklist generator, comparison engine. Interactivity is inherently non-duplicable.
3. **Unique intent** — it serves a genuinely distinct search need, not a synonym of another page.

A town-swapped "Garage Workwear in [Town]" page passes none of these. A "Bulgaria World Cup" page passes none of these *for a workwear site*. Both are retired. This one gate is what converts a 10,000-page liability into a 10,000-page asset.

3. Where genuine uniqueness comes from (the engine)

Stop generating **prose from a template**. Start rendering **data through a component library**. Uniqueness then becomes a by-product of the data, and structure varies because the data varies.

3.1 The data model (entities), not the town list

- **Garments:** hi-vis, coverall, work trouser, FR jacket, safety boot, chef whites, scrub, tabard, softshell, waterproof...
- **Sectors:** construction, healthcare, hospitality, warehousing, automotive, plumbing/electrical, landscaping, FM/cleaning, manufacturing... (*reuse the matrix's taxonomy*)
- **Roles:** welder, electrician, chef, warehouse operative, groundworker, mechanic, nurse...
- **Standards (the richest unique-data seam):** EN ISO 20471 (hi-vis), EN ISO 11612 (heat/flame), EN 343 (rain), EN ISO 20345 (safety footwear), IEC 61482 (arc flash), EN 1149 (antistatic), EN 13034 (chemical), rail GO/RT...
- **Attributes:** fabric, GSM, colour, tape class, features, wash/care temperature, sizing.
- **Intents:** buying guide, "what to wear for X", compliance explainer, care, sizing, comparison, cost, embroidery/branding.

3.2 The supplier-DAM trap (critical — this changes the strategy)

You told me catalogue data and imagery will come from **workwear suppliers' digital asset libraries**. That means the raw specs and photos are **identical to every other reseller** pulling the same DAM — the *opposite* of a unique asset. Rules that follow directly:

- **Never publish supplier copy or specs as page text.** Treat the DAM as *raw input to transform* — synthesise your own spec tables, comparisons, and standard/hazard mappings on top of it.
- **Differentiate imagery.** Raw DAM photos will be image-deduped by Google against dozens of rivals. Plan for own photography where feasible, or at minimum processed/annotated/composited images.
- **The unique value lives in the layer you add** — standards analysis, role/hazard mapping, care guidance, tools, comparisons. No supplier feed contains those. That layer is your moat.

3.3 Discipline against combinatorial near-duplicates

sector × garment × role × town is millions of combinations, ~95% near-duplicate. So:

- **Only build a combination when the data genuinely differs.** "FR coveralls for welders" is real (arc + spatter + radiant heat). "Garage workwear in Abingdon" is not — it's a town swap.
- **Kill the town-multiplication default.** Do **not** build 500 town variants per industry. Build *one strong industry/role guide*, and create a local page **only** where real local data exists (a physical stockist, local regulation, real delivery/lead-time info) — otherwise it's a doorway.
- **Variable depth:** thin nodes become a *section within* a parent page or a filtered view, never a standalone URL.
- **Cap by value, not by a target count.** Retire "build 500" thinking. Build every node that passes §2, and no more.

3.4 Local / geo — the industry-driven model (validated)

Competitors carpet "workwear in [town]" pages, and you asked whether to match them. The answer isn't "no local pages" — it's **no thin ones**. A nationwide shipper's town page can't differ on *service* (same parcel everywhere), so it must differ on **real local-industry data**. The validated model:

- **Drive each town page from its real dominant industries**, not a town-name swap. Aberdeen → offshore (FR/antistatic/multi-norm); Corby → logistics (hi-vis/cold-store); Harrogate → hospitality (uniforms/hygiene). The *workwear answer genuinely differs per town* — that's information gain, and it proved out at **35% shared vocabulary / 4% phrase overlap** across the three prototypes (vs the old site's 77%).
- **An archetype engine, not hand-spinning.** 15 UK economic archetypes each map to an industry mix → standards → garments → Shopify collections; 495 towns are assigned archetype(s), which pre-populate the page. Assignments are confidence-flagged (H/M/L) to **verify against ONS/NOMIS** before build — never invented certainty. (`tranche1-archetypes.csv` , `tranche1-town-list.csv` .)
- **A per-town differentiation stack** (so no single element carries it): real industry mix · sector → standard mapping · a real, anonymised-but-specific **supplied-customer story** (the cleaning site's mechanism — `local-page-data-template.csv` , adapted from the Job_Notes CSV; **real supplies only, never fabricated**) · named local estates/colleges · honest delivery info.
- **Honest schema.** Represent as a delivery/service-area guide (`Article / FAQ / Breadcrumb`), **never a fake LocalBusiness** with a town address you don't have.
- **Depth over breadth + measured.** Build the ~200–500 towns with genuinely rich data (not 5,000); smaller places fold into region pages; run a curated batch, watch it in its own sitemap, then scale.

4. Page taxonomy & template families (fixes the over-templating)

The old build: 2–3 templates for everything. The rebuild: **many families, each structurally distinct because each renders different data**, and each with 3–5 layout variants.

Page type	Example	Why structurally distinct
Standard/regulation explainer	"EN ISO 20471 hi-vis classes explained"	Clause tables, class diagrams, applicability logic
Sector guide (hub)	"Workwear for construction"	Hazard matrix, role list, required standards
Role guide	"What a welder should wear"	Hazard → garment → standard mapping
Garment comparison	"Class 2 vs Class 3 hi-vis"	Decision tables, spec grids from real data
Care & maintenance	"Washing FR coveralls without losing FR"	Step lists, temperature/care tables — high utility
Sizing & fit	Per-garment size guides	Measurement tables, fit diagrams
Interactive tools	Hi-vis class selector; "which coverall?"; PPE checklist; embroidery cost; size calculator	Pure utility — inherently unique
Glossary	Standards/fabric definitions	Short, interlinked, schema-marked
Editorial/blog	Buying guides, deep how-tos	Human-led, real experience (E-E-A-T)

Structural-variation rules in the generator: block order and inclusion driven by *which data fields are present*, so no two pages in a series share an identical skeleton; a component library (spec table, hazard matrix, comparison grid, callout, FAQ, tool embed) assembled by rules; and a QA gate (§6) that rejects any page too similar to its siblings.

Status: every family above except sizing/glossary is now **prototyped** (§0.5) — the sizing, glossary and employer/compliance families reuse the same templates, so the library is effectively complete and the remaining work is content, not design.

4.1 Hub-and-spoke — how the content layer is organised

The useful-content layer (standards explainers, comparisons, care, buying guides ≈ 730 of the ~1,350 core pages) is built as **hub-and-spoke clusters**, mirroring the cleaning site's method:

- A **pillar hub** per topic (e.g. "Hi-vis workwear: classes, standards & how to choose") gives the overview and fans out to the detail.
- **Spokes** each answer one specific question (e.g. "EN ISO 20471 explained", "Class 2 vs Class 3") and link **back to the hub, to sibling spokes, and to the shop collection** as a citation.
- New spokes grow the cluster around the hub over time — the FR cluster, footwear cluster, glove cluster all follow the same recipe. (*Prototyped: the hi-vis cluster — hub + 2 spokes, fully interlinked, 28–41% shared vocab across the three.*)
- Pages are **webpage-styled, not blog-styled** — hero with stat tiles, cards, decision tables, keyfacts, FAQ; no "journal / X-min-read" byline treatment.

5. Branding & design system (fixes problem #1) — systematise the cleaning site

You already have the exemplar. Do for iNeedWorkwear what Phoenix Duct Clean did — **once, centrally** — and render every page through it:

- **One token file:** palette (primary/secondary/neutral/semantic), type scale, spacing, radii, shadows. One palette for the whole site, chosen to sit alongside the future store.
- **Global chrome:** one header, nav, footer, breadcrumb on every page.
- **Component library** themed from tokens (buttons, cards, tables, callouts, tool shells).
- **Imagery system** (see §3.2 — differentiate from DAM stock).
- **Accessibility + Core Web Vitals** budgets baked in (both are quality signals).
- **Deliverable:** a brand spec + a live style-guide page, before any content page is generated.

Status: the iNeedWorkwear system is **established and in use across all prototypes** — deep navy ground, workwear-orange primary, hi-vis-lime interactive accent; Barlow Condensed (display) / Inter (body) / IBM Plex Mono (labels); one header/footer/nav. Deliberately its own brand, built with the cleaning site's discipline. Next step is extracting it into a reusable tokens + component starter.

6. Build system & QA gates (what makes scale safe)

6.1 Stack [to confirm]

- **Static-site generator** (Astro or Next.js) — component-driven, fast, indexable.
- **Content as data:** the entity dataset in structured files/DB, rendered via template families and the component library; editorial in MDX/CMS. Copy is never hand-duplicated across pages.
- **Clean separation:** `data/` (facts) · `templates/` (families) · `components/` (blocks) · `content/` (editorial).

6.2 Pre-publish gates (automated in CI — nothing ships that fails)

1. **Similarity gate** — embeddings or SimHash/MinHash vs siblings; reject above a threshold. (*This alone would have blocked the `am-aberdeen` / `am-abingdon` pattern.*)
2. **Thin-content gate** — must contain real data blocks (spec table, hazard matrix, tool), not just prose.
3. **§2 value-test gate** — each page tagged with which test it passes; untagged = blocked.
4. **Doorway gate** — CTA/shop-link-to-content ratio ceiling. (*This would have blocked the World Cup pages.*)
5. **Topical-relevance gate** — page must map to a workwear entity; off-topic = blocked.
6. **Technical SEO** — unique title/meta/H1, correct canonical, schema.org, internal links, no orphans.
7. **Design/lint** — design-system + accessibility + CWV budgets.

6.3 Index management

- **Phased publishing** in tranches; watch indexation and rankings before the next tranche.
- `noindex` until proven for experimental families.

- **Sitemaps segmented by page type** to monitor how Google treats each family.

6.4 Staging & safe launch (build in the open, invisible to Google)

You can build the whole site publicly and preview it as you go — the trick is *how* you hide it.

Don't rely on `robots.txt` `Disallow` alone. It blocks crawling, not indexing: a disallowed URL can still be indexed as a bare link, and — critically — if a page is disallowed, Google can't read a `noindex` tag on it, so the two cancel out. Use one of the two correct methods instead:

Method	Publicly reachable?	In Google's index?	Use when
Password (HTTP auth) — cPanel → <i>Directory Privacy</i>	No (login required)	Never	Private team/stakeholder preview — recommended
Site-wide <code>noindex</code> — X-Robots-Tag: <code>noindex</code> header or <code><meta name="robots" content="noindex"></code> in every page head, <code>robots.txt</code> left open so Google can read it	Yes	No	You want a shareable public URL

- **cPanel note:** password protection is **Directory Privacy** (a login wall), *not* file-permission codes (644 / 755). Leave permissions at the web defaults — 700 on a folder just breaks the site.
- **Build on a staging subdomain/domain** (e.g. `staging.ineedworkwear.co.uk`), not the live path, so no half-built signals ever touch the production domain.
- **Go-live checklist (per tranche):** the #1 SEO disaster is launching with `noindex /auth` still on. Before lifting the block for a tranche: content live **and** matching Shopify collections live · no dead links · titles/meta/canonical/schema correct · then remove the block, submit the segmented sitemap in Search Console, and confirm indexing.

6.5 The canary gate — decide *when* to scale (the current blocker)

The cleaning site launched only days ago, so whether the programmatic model passes *Google's* quality bar (not just ours) is unknown — and takes weeks to read. Rather than bet the workwear domain and the future store on a guess, **use the cleaning 3,000 as the canary**: read Search Console before replicating the model on workwear. Full method in `canary-watch-framework.md` ; in brief:

- **Earliest, most honest signal = indexation ratio** (GSC → Pages): of the URLs submitted, how many index vs sit in "*Crawled / Discovered – currently not indexed*"? A large, static pile there is Google declining pages it has seen.
- **Segment the sitemap by page type** so indexation is read *per pattern* — the key diagnostic.
- **Go/Amber/Red thresholds:** 🟢 ≥70–80% indexed and climbing by wk 6–8 → replicate. 🟡 40–70% or one weak family → fix that pattern, re-test, replicate only proven families. 🔴 <40% flat, duplicate-canonical statuses, or spike-then-decay → don't touch workwear; rebuild the model first.
- **Durability is the real test** — hold at wk 12–16 and across the first core update; an early bump is a "honeymoon", not proof. Keep the workwear build ready but **unpublished** until green.

(Assets: `canary-watch-framework.md`, `canary-tracking-sheet.csv`, `canary-watch-dashboard.html`.)

7. Dealing with the existing pages

They are actively dragging any shared quality signal. Triage every URL:

1. **Off-topic traffic-net series** (world cup, festivals, eat, youth clubs, restaurants, nurseries, town criers): **delete** → **410** (or redirect to the closest relevant hub only where one exists). Do not migrate. These are pure liability.
2. **Town-swapped industry pages: retire the town variants.** Keep the *industry topic*; rebuild as a smaller set of deep sector/role guides on the new system. 301 each dead town URL to its sector hub.
3. **Genuine editorial** (real buying guides): **rebuild** as deep, first-hand editorial. Keep the URL where it has any equity.
4. **Migration map** with 301s; no soft-404s left behind.

Expect to delete the large majority of current URLs. Removing thin/off-topic pages *raises* the site's average quality — it's a recovery lever, not a loss.

8. Cross-site linking (protect both domains)

- **Every link to the store is a relevant citation** ("coveralls certified to EN ISO 11612 → view specs"), never a blanket funnel CTA. Relevance is what separates a resource from a doorway.
 - **The content site must stand alone** — useful with zero purchases. If a page passes §2 on its own merits, it is not a doorway by definition.
 - **Content first, commercial links secondary** (the opposite of the current World Cup pages).
 - Because the store isn't built yet, **design both around one brand system now** so they read as one family when the store launches.
-

9. Metrics & governance

- **Value density** — % of indexed pages passing §2 (north star; target 100%).
 - **Indexation ratio** per page-type sitemap (near-duplicate families show low ratios).
 - **Internal similarity distribution** — duplicate-cluster report over time.
 - **Organic performance** per family + funnel referrals to the store.
 - **CWV + design conformance** — automated budgets.
 - **Editorial calendar** for the human-led layer (freshness + E-E-A-T).
-

10. Phased roadmap

Phase	Outcome	Status
A. Strategy & evidence	Diagnosis + governing rules	✅ Done (\$0–\$3)
B. Design system	"One site" identity	✅ Established (\$5, \$0.5)
C. Template library	A prototype per page type	✅ Done — role/local/hub/spoke/care (\$0.5)
D. Page architecture & data engine	Counts + local model + town data	✅ Done — 14 families, archetype engine, 495 towns (\$0.5)
E. Canary gate	Decide <i>when</i> to scale	🟡 In progress — the current blocker (\$6.5)
0. Decommission risk	Stop the bleeding	❑ Delete/ 410 off-topic series; 301 town pages to hubs; migration map (\$7)
1. Cluster spec + build harness	Repeatable build	❑ "Hub + N spokes" recipe + data schema; QA gates in CI (\$6.2); tokens/component starter
2. Data model → catalogue	Facts, not prose	❑ Entity dataset; DAM-transform pipeline; Shopify collection + tag structure (per tranche)
3. Pilot tranche	Prove on workwear	❑ 1–2 families + matching collections, staged, indexed, monitored
4. Scale out	Thousands, safely	❑ Remaining families/clusters in monitored tranches
5. Tools + editorial	Utility + E-E-A-T moat	❑ Flagship tools; ongoing editorial

Phases A–D are complete; **E (canary) is the gate**. Phases 0–5 fire once E reads green — 0 (cleanup) and 1 (harness) can run in parallel with the canary watch since neither publishes live workwear content.

11. Open questions / decisions

Resolved this session: **local/geo** (industry-driven model, \$3.4) · **content sequencing** (link to Shopify collections, not products — many-to-few, category-parallel tranches) · **staging** (Directory Privacy, \$6.4). Still open:

1. **Scope of the cull** — delete the off-topic series outright (my strong recommendation), or a page-by-page review first?
2. **Stack** — Astro/Next, or stay on the current cPanel/hand-built hosting? (The prototypes are self-contained HTML that already run on cPanel, so either is viable.)
3. **Imagery** — any own photography, or DAM-only at launch? (Sets how hard we work on image differentiation.)
4. **Are the cleaning site's stats/testimonials real?** (Determines whether that pattern is safe to reuse.)
5. **URL/domain** — same `iNeedWorkwear` domain with redirects, or fresh structure?
6. **Do you have real workwear supply records** (like the cleaning Job_Notes CSV) to seed the supplied-customer stories, or does that layer build up as you win customers?

Appendix A — Worked example: one node, done right

Topic: "Flame-resistant coveralls for welders"

- **Passes §2:** unique information (welding hazards → EN ISO 11612 + IEC 61482) **and** unique utility (embedded FR-standard finder).
- **Data rendered:** required-standards table, fabric/GSM options (transformed from DAM, not pasted), care rules that preserve FR properties, layering, sizing.
- **Structure:** hazard matrix → standards table → spec comparison → care do/don't → tool → FAQ. A different DOM from a sizing page, because it contains different data.
- **Outbound link:** contextual — "coveralls certified to EN ISO 11612 →" (citation, not funnel).
- **Result:** useful with zero purchase; not thin, not duplicate, not a doorway.

Contrast with the retired pattern: `am-abingdon` — "Garage and Mechanic Workwear Supplier" with the town swapped and 7/8 headings identical to `am-aberdeen`; or `world-cup-bulgaria` — a football-history essay stuffed with 31 "workwear" mentions and 11 shop links. Both fail all three tests. **That is the build we are replacing.**

Appendix B — Asset index (everything built this project)

All committed to branch `claude/workwear-restructure-plan-mq8w1e`.

Strategy & planning (docs/)

- `workwear-restructure-plan.md / .pdf` — this plan
- `page-architecture.csv` — 14 page families, counts, tiers, value tests, collections
- `page-architecture-entities.csv` — the standards/sectors/roles/garments axes behind the counts
- `page-architecture-matrix.html` — visual page-architecture dashboard
- `local-page-data-template.csv` — the per-town input schema (Aberdeen + Corby worked examples)
- `tranche1-archetypes.csv` — 15 UK economic archetypes → industry mix/standards/garments/collections
- `tranche1-town-list.csv` — 495 towns, archetype-assigned, confidence-flagged
- `tranche1-generator.mjs` — regenerates the town list from the archetypes
- `canary-watch-framework.md` — the pre-scale go/no-go gate
- `canary-tracking-sheet.csv` — weekly Search Console log template
- `canary-watch-dashboard.html` — visual canary go/no-go dashboard

Prototypes (prototypes/) — one per distinct page type

- `fr-coveralls-for-welders.html` — role guide + interactive standard-finder tool
- `workwear-aberdeen.html` · `workwear-corby.html` · `workwear-harrogate.html` — industry-driven local pages
- `hi-vis-workwear-guide.html` — content **hub** (pillar)
- `en-iso-20471-hi-vis-classes-explained.html` — **spoke** (standard explainer)

- `class-2-vs-class-3-hi-vis.html` — **spoke** (comparison / decision)
- `how-to-wash-fr-coveralls.html` — care / maintenance (HowTo)

Immediate next steps (order): finish the **canary gate** (§6.5) → in parallel, **decommission risk** (§7) and author the **cluster spec + build harness** (Phase 1) → on green, **pilot tranche**.